

MODEL UNCERTAINTY IN CROSS-COUNTRY GROWTH REGRESSIONS

CARMEN FERNÁNDEZ,^a EDUARDO LEY^{b*} AND MARK F. J. STEEL^c

^a *School of Mathematics and Statistics, University of Saint Andrews, UK*

^b *IMF Institute, International Monetary Fund, Washington, DC, USA*

^c *Institute of Mathematics and Statistics, University of Kent at Canterbury, UK*

SUMMARY

We investigate the issue of model uncertainty in cross-country growth regressions using Bayesian Model Averaging (BMA). We find that the posterior probability is spread widely among many models, suggesting the superiority of BMA over choosing any single model. Out-of-sample predictive results support this claim. In contrast to Levine and Renelt (1992), our results broadly support the more ‘optimistic’ conclusion of Sala-i-Martin (1997b), namely that some variables are important regressors for explaining cross-country growth patterns. However, care should be taken in the methodology employed. The approach proposed here is firmly grounded in statistical theory and immediately leads to posterior and predictive inference. Copyright © 2001 John Wiley & Sons, Ltd.

1. INTRODUCTION

Many empirical studies of the growth of countries attempt to identify the factors explaining the differences in growth rates by regressing observed GDP growth on a host of country characteristics that could possibly affect growth. This line of research was heavily influenced by Kormendi and Meguire (1985) and Barro (1991). Excellent recent surveys of these cross-section studies and their role in the broader context of economic growth theory are provided in Durlauf and Quah (1999) and Temple (1999). A more specific discussion of various approaches to model uncertainty in this context can be found in Temple (2000) and Brock and Durlauf (2000). The latter paper advocates a decision-theoretic approach to policy-relevant empirical analysis.

In this paper we focus on cross-country growth regressions and attempt to shed further light on the importance of such models for empirical growth research. Prompted by the proliferation of possible explanatory variables in such regressions and the relative absence of guidance from economic theory as to which variables to include, Levine and Renelt (1992) investigate the ‘robustness’ of the results from such linear regression models. They use a variant of the Extreme-Bounds Analysis introduced in Leamer (1983, 1985) and conclude that very few regressors pass the extreme-bounds test. In response to this rather negative finding, Sala-i-Martin (1997b) employs a less severe test for the importance of explanatory variables in growth regressions, the aim being ‘to assign some level of confidence to each of the variables’¹ rather than to classify them as robust

* Correspondence to: Eduardo Ley, International Monetary Fund, 700 19 Street NW, Washington DC 20431, USA.

¹ For each variable, he denotes the level of confidence by CDF(0) and defines it as the maximum of the probability mass to the left and the right of zero for a (Normal) distribution centred at the estimated value of the regression coefficient and with the corresponding estimated variance. He deals with model uncertainty by running many different regressions and either computing CDF(0) based on the averages of the estimated means and variances (approach 1), or redefining CDF(0)

versus non-robust. On the basis of his methodology, Sala-i-Martin (1997b) identifies a relatively large number of variables as important for growth regression.

Here we set out to investigate this issue in a formal statistical framework that explicitly allows for the specification uncertainty described above. In particular, a Bayesian framework allows us to deal with both model and parameter uncertainty in a straightforward and formal way. We also consider an extremely large set of possible models by allowing for any subset of up to 41 regressors to be included in the model. This means we have a set of $2^{41} = 2.2 \times 10^{12}$ (over two trillion!) different models to deal with. Novel Markov chain Monte Carlo (MCMC) techniques are adopted to solve this numerical problem, using the so-called Markov chain Monte Carlo Model Composition (MC³) sampler, first used in Madigan and York (1995).

Our findings are based on the same data as those of Sala-i-Martin² and broadly support the more 'optimistic' conclusion of Sala-i-Martin (1997b), namely that some variables are important regressors for explaining cross-country growth patterns. However, the variables we identify as most useful for growth regression differ somewhat from his results. More importantly, we do not advocate selecting a subset of the regressors, but we use Bayesian Model Averaging, where all inference is averaged over models, using the corresponding posterior model probabilities as weights. It is important to point out that our methodology allows us to go substantially further than the previous studies, in that we provide a clear interpretation of our results and a formal statistical basis for inference on parameters and out-of-sample prediction. Finally, let us briefly mention that this paper is solely intended to investigate a novel methodology to tackle the issues of model uncertainty and inference in cross-country growth regressions, based on the Normal linear model. We do not attempt to address here the myriad other interesting topics, such as convergence of countries, data quality or any further issues of model specification.

2. THE MODEL AND THE METHODOLOGY

Following the analyses in Levine and Renelt (1992) and Sala-i-Martin (1997b) as well as the tradition in the growth regression literature, we will consider linear regression models where GDP growth for n countries, grouped in a vector y , is regressed on an intercept, say α , and a number of explanatory variables chosen from a set of k variables in a matrix Z of dimension $n \times k$. Throughout, we assume that $\text{rank}(\iota_n : Z) = k + 1$, where ι_n is an n -dimensional vector of 1's, and define β as the full k -dimensional vector of regression coefficients.

Whereas Levine and Renelt and Sala-i-Martin restrict the set of regressors to always contain certain key variables and then allow for four³ other variables to be added, we shall allow for any subset of the variables in Z to appear in the model. This results in 2^k possible models, which will thus be characterized by the selection of regressors. We denote by M_j the model with regressors grouped in Z_j , leading to

$$y = \alpha \iota_n + Z_j \beta_j + \sigma \varepsilon \quad (1)$$

as the average of the CDF(0)'s resulting from the various regressions (approach 2). In both cases, the averaging over models is either done uniformly or with weights proportional to the likelihoods. See also our footnote 14 in this context. Regressors leading to $\text{CDF}(0) > 0.95$ are classified as 'significant'.

² We thank Xavier Sala-i-Martin for making his data publicly available at his website. The data are also available at this journal's website.

³ Levine and Renelt (1992) consider one up to four added regressors, Sala-i-Martin (1997a,b) restricts the analysis to exactly four extra regressors.

where $\beta_j \in \Re^{k_j} (0 \leq k_j \leq k)$ groups the relevant regression coefficients and $\sigma \in \Re_+$ is a scale parameter. In line with most of the literature in this area (see e.g. Mitchell and Beauchamp, 1988, and Raftery, Madigan and Hoeting, 1997), exclusion of a regressor means that the corresponding element of β is zero. Thus, we are always conditioning on the full set of regressors Z . Finally, we shall assume that ε follows an n -dimensional Normal distribution with zero mean and identity covariance matrix.

In our Bayesian framework, we need to complete the above sampling model with a prior distribution for the parameters in M_j , namely α , β_j and σ . In the context of model uncertainty, it is acknowledged that the choice of this distribution can have a substantial impact on posterior model probabilities (see e.g. Kass and Raftery, 1995, and George, 1999). Raftery *et al.* (1997) use a ‘weakly-informative’ prior which is data-dependent. Here we follow Fernández, Ley and Steel (2001) who, on the basis of theoretical results and extensive simulations, propose a ‘benchmark’ prior distribution that has little influence on posterior inference and predictive results and is, thus, recommended for the common situation in which incorporating substantive prior information into the analysis is not possible or desirable. In particular, they propose to use improper noninformative priors for the parameters that are common to all models, namely α and σ , and a g -prior structure for β_j . This corresponds to the product of

$$p(\alpha, \sigma) \propto \sigma^{-1} \quad (2)$$

and

$$p(\beta_j | \alpha, \sigma, M_j) = f_N^{k_j}(\beta_j | 0, \sigma^2 (g Z_j' Z_j)^{-1}) \quad (3)$$

where $f_N^q(w|m, V)$ denotes the density function of a q -dimensional Normal distribution on w with mean m and covariance matrix V . Fernández *et al.* (2001) investigate many possible choices for g in equation (3) and conclude that taking $g = 1/\max\{n, k^2\}$ leads to reasonable results. Finally, the $k - k_j$ components of β which do not appear in M_j are exactly equal to zero. Note that in equation (2) we have assumed a common prior for σ across the different models. This is a usual practice in the literature (e.g. Mitchell and Beauchamp, 1988; Raftery *et al.*, 1997) and does not seem unreasonable since, by always conditioning on the full set of regressors, σ keeps the same meaning (namely the residual standard deviation of y given Z) across models. The distribution in equation (2) is the standard non-informative prior for location and scale parameters, and is the only one that is invariant under location and scale transformations (such as induced by a change in the units of measurement). As Fernández *et al.* (2001) show, the prior in equations (2) and (3) has convenient properties (marginal likelihoods can be computed analytically) while leading to satisfactory results from a posterior and predictive point of view.

So far, we have described the sampling and prior setting under model M_j . As already mentioned, a key aspect of the problem is the uncertainty about the choice of regressors—i.e. model uncertainty. This means that we also need to specify a prior distribution over the space \mathcal{M} of all 2^k possible models:

$$P(M_j) = p_j, \quad j = 1, \dots, 2^k, \quad \text{with } p_j > 0 \text{ and } \sum_{j=1}^{2^k} p_j = 1 \quad (4)$$

In our empirical application, we will take $p_j = 2^{-k}$ so that we have a Uniform distribution on the model space. This implies that the prior probability of including a regressor is 1/2, independently of the other regressors included in the model. This is a standard choice in the absence of prior

information but other choices—e.g. downweighing models with a large number of regressors—are certainly possible. See Chipman (1996) for priors that allow for dependence between regressors.

The Bayesian paradigm now automatically deals with model uncertainty, since the posterior distribution of any quantity of interest, say Δ , is an average of the posterior distributions of that quantity under each of the models with weights given by the posterior model probabilities. Thus

$$P_{\Delta|y} = \sum_{j=1}^{2^k} P_{\Delta|y, M_j} P(M_j|y) \quad (5)$$

Note that, by making appropriate choices of Δ , this formula gives the posterior distribution of parameters such as the regression coefficients or the predictive distribution that allows to forecast future or missing observables. The marginal posterior probability of including a certain variable is simply the sum of the posterior probabilities of all models that contain this regressor. The procedure described in equation (5), which is typically referred to as Bayesian Model Averaging (BMA), immediately follows from the rules of probability theory—see e.g. Leamer (1978).

We now turn to the issue of how to compute $P_{\Delta|y}$ in equation (5). The posterior distribution of Δ under model M_j , $P_{\Delta|y, M_j}$, is typically of standard form (the following sections mention this distribution for several choices of Δ). The additional burden due to model uncertainty is having to compute the posterior model probabilities, which are given by

$$P(M_j|y) = \frac{l_y(M_j)p_j}{\sum_{h=1}^{2^k} l_y(M_h)p_h} \quad (6)$$

where $l_y(M_j)$, the marginal likelihood of model M_j , is obtained as

$$l_y(M_j) = \int p(y|\alpha, \beta_j, \sigma, M_j) p(\alpha, \sigma) p(\beta_j|\alpha, \sigma, M_j) d\alpha d\beta_j d\sigma \quad (7)$$

with $p(y|\alpha, \beta_j, \sigma, M_j)$ the sampling model corresponding to equation (1) and $p(\alpha, \sigma)$ and $p(\beta_j|\alpha, \sigma, M_j)$ the priors defined in equations (2) and (3), respectively. Fernández *et al.* (2001) show that for the Bayesian model in (1)–(4) the marginal likelihood can be computed analytically. In the somewhat simplifying case where, without loss of generality, the regressors are demeaned, such that $\iota'_n Z = 0$, and defining $X_j = (\iota_n : Z_j)$, $\bar{y} = \iota'_n y/n$ and $M_{X_j} = I_n - X_j(X'_j X_j)^{-1} X'_j$, they obtain

$$l_y(M_j) \propto \left(\frac{g}{g+1} \right)^{k_j/2} \left(\frac{1}{g+1} y' M_{X_j} y + \frac{g}{g+1} (y - \bar{y} \iota_n)' (y - \bar{y} \iota_n) \right)^{-(n-1)/2} \quad (8)$$

Since marginal likelihoods can be computed analytically, the same holds for the posterior model probabilities, given in equation (6), and the distribution described in equation (5).

In practice, however, computing the relevant posterior or predictive distribution through equations (5), (6) and (8) is hampered by the very large amount of terms involved in the sums. In our application, we have $k = 41$ possible regressors, and we would thus need to calculate posterior probabilities for each of the $2^{41} = 2.2 \times 10^{12}$ models and average the required distributions over all these models. Exhaustive evaluation of all these terms is computationally prohibitive. Even using fast updating schemes, such as that proposed by Smith and Kohn (1996), in combination with the

Gray code order, computations become practically infeasible when k is larger than approximately 25 (George and McCulloch, 1997).⁴ In order to substantially reduce the computational effort, we shall approximate the posterior distribution on the model space \mathcal{M} by simulating a sample from it, applying the MC³ methodology of Madigan and York (1995). This consists in a Metropolis algorithm—see e.g. Tierney (1994) and Chib and Greenberg (1996)—to generate drawings through a Markov chain on \mathcal{M} which has the posterior model distribution as its stationary distribution. The sampler works as follows. Given that the chain is currently at model M_s , a new model M_j is proposed randomly through a Uniform distribution on the space containing M_s , and all models with either one regressor more or one regressor less than M_s . The chain moves to M_j with probability $p = \min\{1, [l_y(M_j)p_j]/[l_y(M_s)p_s]\}$ and remains at M_s with probability $1 - p$. Raftery *et al.* (1997) and Fernández *et al.* (2001) use MC³ methods in the context of the linear regression model.

In the implementation of MC³, we shall take advantage of the fact that marginal likelihoods can be computed analytically through equation (8). Thus, we shall use the chain to merely indicate which models should be taken into account in computing the sums in equations (5) and (6)—i.e. to identify the models with high posterior probability. For the set of models visited by the chain, posterior probabilities will be computed through appropriate normalization of $l_y(M_j)p_j$. This idea was called ‘window estimation’ in Clyde, Desimone and Parmigiani (1996) and was denoted by ‘Bayesian Random Search’ in Lee (1996). In addition, Fernández *et al.* (2001) propose to use this as a convenient diagnostic aid for assessing the performance of the chain. A high positive correlation between posterior model probabilities based on the empirical frequencies of visits in the chain, on the one hand, and the exact marginal likelihoods, on the other, suggests that the chain has reached its equilibrium distribution. Of course, the chain will not cover the entire model space since this would require sampling all 2^{41} models, an impossible task as we already mentioned. Thus, the sample will not constitute, as such, a perfect replica of the posterior model distribution. Rather, the objective of sampling methods in this context is to explore the model space in order to capture the models with higher posterior probability (George and McCulloch, 1997). Nevertheless, our efficient implementation (see footnote 13) allows us to cover a high percentage of the posterior mass and, thus, to also characterize a very substantial amount of the variability inherent in the posterior distribution.

3. POSTERIOR RESULTS

We take the same data as used and described in Sala-i-Martin (1997b), covering 140 countries, for which average per capita GDP growth was computed over the period 1960–1992. Sala-i-Martin starts with the model in (1) and a large set of 62 variables that could serve as regressors.⁵ He then

⁴ The largest computational burden lies on the evaluation of the marginal likelihood of each model. In the context of our application and without use of fast updating schemes or Gray code ordering, an estimate of the average rate at which we can compute $l_y(M_j)$ is over 36,000 per minute of CPU time (on a state-of-the-art Sun Ultra-2 with two 296MHz CPUs, 512Mb of RAM and 3.0Gb of swap space running under Solaris 2.6), which would imply that exhaustive evaluation would approximately take 115 years. Even if fast updating algorithms and ordering schemes can be found that reduce this by a factor $k = 41$ (as suggested in George and McCulloch, 1997), this is still prohibitive. In addition, storing and manipulating the results—e.g. to compute predictive and posterior densities for regression coefficients—seems totally unfeasible with current computing technology.

⁵ This set of possible regressors does not include the investment share of GDP, which was included in Levine and Renelt (1992) as a variable always retained in the regressions. However, they comment that then the only channel through

restricts his analysis to those models where three specific variables are always included (these are the level of GDP, life expectancy and primary school enrollment rate, all for 1960) and, for each of the remaining 59 variables, he adds that variable and all possible triplets of the other 58. He finally computes⁶ CDF(0) for that regressor as the weighted average of the resulting CDF(0)'s to conclude that 22 of the 59 variables are 'significant', in that CDF(0) is larger than 0.95. Thus, in all he considers 455,126 different models,⁷ which we will denote by \mathcal{M}_s .

We shall undertake our analysis on the basis of the following set of regressors. First, we take the 25 variables that Sala-i-Martin (1997b) flagged as being important (his three retained variables and the 22 variables in his Table 1, page 181). We have available $n = 72$ observations for all these regressors. We then add to this set all regressors that do not entail a loss in the number of observations. Thus, we keep $n = 72$ observations which allows us to expand the set of regressors to a total of $k = 41$ possible variables. Z will be the 72×41 design matrix corresponding to these variables (transformed by subtracting the mean, so that $\iota_n'Z = 0$), and we shall allow for any subset of these 41 regressors, leading to a total set of $2^{41} = 2.2 \times 10^{12}$ models under consideration in \mathcal{M} . Since we do not start from the full set of 62 variables, we do not cover all models in \mathcal{M}_s . On the other hand, since we allow for any combination of regressors, our model space is of much larger size than \mathcal{M}_s . Clearly, we cover the subset of \mathcal{M}_s that corresponds to the 41 regressors considered here. This intersection between \mathcal{M} and \mathcal{M}_s consists of 73,815 models and will be denoted by \mathcal{M}_I in the sequel. In view of the fact that \mathcal{M}_I contains all models in \mathcal{M}_s using Sala-i-Martin's (1997b) favoured regressors, we would certainly expect that a relatively large fraction of the posterior mass in \mathcal{M}_s is concentrated in \mathcal{M}_I .

To analyse these data, we use the Bayesian model in equations (1)–(4) with a Uniform prior on model probabilities, i.e. $p_j = 2^{-k}$ in (4). Since $n < k^2$, we shall take $g = 1/k^2$ in the prior in (3). Given the size of \mathcal{M} we would expect to need a fairly large amount of drawings of the MC³ sampler to adequately identify the high posterior probability models. We shall report results from a run with 2 million recorded drawings after a burn-in of 1 million discarded drawings, leading to a correlation coefficient between visit frequencies and posterior probabilities based on (8) of 0.993. The results based on a different run with 500,000 drawings after a mere 100,000 burn-in drawings are very close indeed. In particular, the best 76 models (those with posterior mass above 0.1%) are exactly the same in both runs. Many more runs, started from randomly drawn points in model space and leading to virtually identical results, confirmed the good behaviour of the sampler. More formally, we can estimate the total posterior model probability visited by the chain following George and McCulloch (1997), by comparing visit frequencies and the aggregate marginal likelihood for a predetermined subset of models. Basing such an estimate on the best models visited in a short run, we estimate the model probability covered by the reported chain to be 70%, which is quite reasonable in view of the fact that we only visit about one in every 15 million models.⁸

which other variables can affect growth is the efficiency of resource allocation. Sala-i-Martin (1997a) finds that including investment share in all regressions does not critically alter the conclusions with respect to those of Sala-i-Martin (1997b).

⁶ Approach 2 described in his footnote 1. He comments that levels of significance found using approach 1 in footnote 1 were virtually identical. However, see our discussion at the end of this section.

⁷ Note that the (almost) 2 million regressions in the title of his paper result from counting each model four times, since he distinguishes between identical models according to whether a variable is being "tested" or merely added in the triplet (see also his footnote 3).

⁸ In order to put this covered probability estimate in perspective, we have to consider the trade-off between accuracy and computing effort. If we run a longer chain of 5 million draws, after a 1 million burn-in, we visit about twice as many

Table I. Marginal evidence of importance

	Regressors	BMA Post.prob.	Sala-i-Martin CDF(0)
⇒	1 GDP level in 1960	1.000	1.000
→	2 Fraction Confucian	0.995	1.000
⇒	3 Life expectancy	0.946	0.999
→	4 Equipment investment	0.942	1.000
→	5 Sub-Saharan dummy	0.757	0.997
→	6 Fraction Muslim	0.656	1.000
→	7 Rule of law	0.516	1.000
→	8 Number of Years open economy	0.502	1.000
→	9 Degree of Capitalism	0.471	0.987
→	10 Fraction Protestant	0.461	0.966
→	11 Fraction GDP in mining	0.441	0.994
→	12 Non-Equipment Investment	0.431	0.982
→	13 Latin American dummy	0.190	0.998
⇒	14 Primary School Enrollment, 1960	0.184	0.992
→	15 Fraction Buddhist	0.167	0.964
→	16 Black Market Premium	0.157	0.825
→	17 Fraction Catholic	0.110	0.963
→	18 Civil Liberties	0.100	0.997
	19 Fraction Hindu	0.097	0.654
→	20 Primary exports, 1970	0.071	0.990
→	21 Political Rights	0.069	0.998
→	22 Exchange rate distortions	0.060	0.968
	23 Age	0.058	0.903
→	24 War dummy	0.052	0.984
	25 Size labor force	0.047	0.835
	26 Fraction speaking foreign language	0.047	0.831
	27 Fraction of Pop. Speaking English	0.047	0.910
	28 Ethnolinguistic fractionalization	0.035	0.643
→	29 Spanish Colony dummy	0.034	0.938
→	30 SD of black-market premium	0.031	0.993
	31 French Colony dummy	0.031	0.702
→	32 Absolute latitude	0.024	0.980
	33 Ratio workers to population	0.024	0.766
	34 Higher education enrollment	0.024	0.579
	35 Population Growth	0.022	0.807
	36 British Colony dummy	0.022	0.579
	37 Outward Orientation	0.021	0.634
	38 Fraction Jewish	0.019	0.747
→	39 Revolutions and coups	0.017	0.995
	40 Public Education Share	0.016	0.580
	41 Area (Scale Effect)	0.016	0.532

Let us now focus on the results obtained in a typical chain of 2 million recorded draws: 149,507 models are visited and the best model obtains a posterior probability of 1.24%. The mass is very spread out: the best 25,219 models only cover 90% of the posterior model probability, making BMA a necessity for meaningful inference. The 76 models with posterior probabilities over 0.1% all have in between 6 and 12 regressors,⁹ which is not greatly at odds with the 7-regressor models

models, which account for an estimated 77% of the posterior mass, at the cost of more than tripling the CPU requirements. Posterior results are virtually unaffected as all added models have very low posterior probability.

⁹ The Bayes factor, obtained from equation (8) as $B_{js} = l_y(M_j)/l_y(M_s)$, has a built-in mechanism to avoid overfitting through the first factor. It can easily be seen that for the values of g and n used here, having one more regressor should be

in \mathcal{M}_S . Indeed, \mathcal{M}_I , the intersection of \mathcal{M} and \mathcal{M}_S , is allocated 0.38% posterior probability, which is 112,832 times the prior mass.¹⁰ So there is a small but non-negligible amount of support for the class of models chosen by Sala-i-Martin, and the best model in \mathcal{M}_I receives a posterior probability of 0.30%.

Marginal posterior probabilities of including each of the 41 regressors $CDF(0)$ that were at the basis of the findings of Sala-i-Martin (1997b). An arrow in front of a regressor identifies the 22 important regressors of Sala-i-Martin (1997b, Table 1) and regressors with double arrows are the ones he always retained in the models. Starting with the latter three, it is clear that GDP in 1960 and (to a lesser degree) life expectancy can indeed be retained without many problems, but that is not the case for primary school enrollment. The 22 regressors that Sala-i-Martin flags as important have posterior probabilities of inclusion ranging from as low as 1.7% to 100.0%. Nevertheless, the Spearman rank correlation coefficient between $CDF(0)$ and marginal posterior inclusion probabilities is 0.94.

Of course, this is only a small part of the information provided by BMA, which really provides a joint posterior distribution of all possible 41 regression coefficients, consisting of both discrete (at zero, when a regressor is excluded) and continuous parts. Among other things, it will provide information on which combinations of regressors are likely to occur, avoiding models with highly collinear regressors. For example, Civil liberties and Political rights are the two variables with the highest pairwise correlation in the sample: it is 0.96. Thus, it is quite likely that if one of these variables is included in a model, the other will not be needed, as it captures more or less the same information. Indeed, the posterior probability of including both these variables is 0.1%, which is much smaller than the product of the marginal inclusion probabilities (about 0.7%).

Thus, the marginal importance of regressors derived from our methodology does not lead to the same results as found in Sala-i-Martin (1997b), but the results are not too dissimilar either. However, there are a number of crucial differences. First, BMA addresses the issue of probabilities of models, not just of individual regressors, and thus provides a much richer type of information than simply that indicated by Table I. In addition, an important difference is that we have a coherent statistical framework in which inference can be based on all models, averaged with their posterior probabilities. So there is no need for us to choose a particular model or discard any of the regressors. As we shall see in the next section, using BMA rather than choosing one particular model is quite beneficial for prediction. In contrast, there is no formal inferential or modelling recipe attached to the conclusions in Sala-i-Martin. Does one adopt a model with all important regressors included at the same time or with a subset of four of those? If model averaging is the implicit message, it is unclear how to implement this outside the Bayesian paradigm.

Figure 1 graphically presents the marginal posterior distribution of some regression coefficients. A gauge on top of the graphs indicates (in black) the posterior probability of inclusion of the corresponding regressor (thus revealing the same information as in Table I). The density in each of the graphs describes the posterior distribution of the regression coefficient given that the corresponding variable is included in the regression. Each of these densities is itself a mixture as in equation (5) over the Student- t posteriors for each model that includes that particular regressor. A dashed vertical line indicates the averaged point estimate presented in Table I of

offset by a decrease in the sum of squared residuals of about 10% in order to have a unitary Bayes factor. For asymptotic links between our Bayes factors and various classical model selection criteria, see Fernández *et al.* (2001).

¹⁰ For comparison, the posterior probability of the best 73,815 models (the same number of models as in \mathcal{M}_I) multiplies the corresponding prior probability by about 29.5 million.

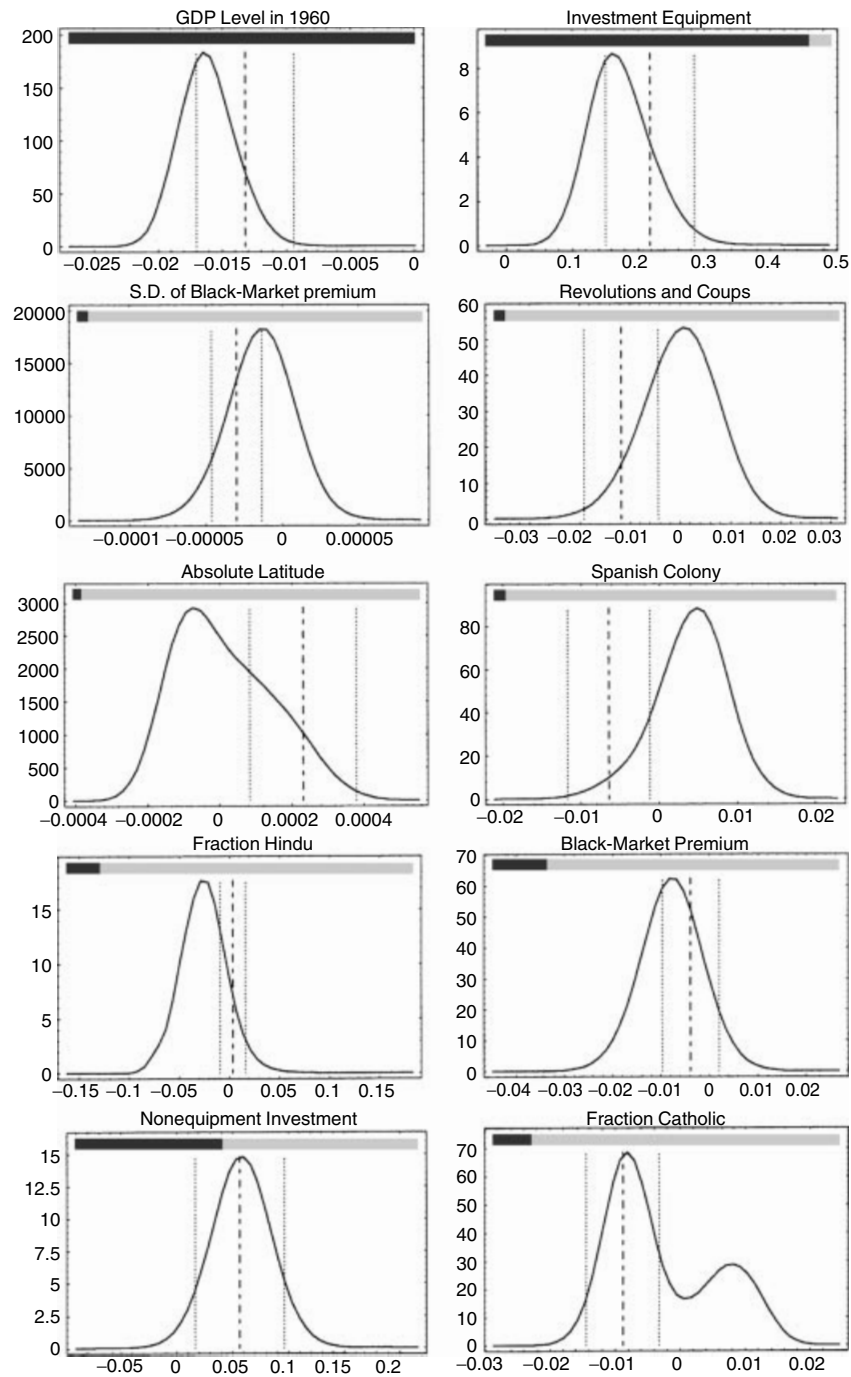


Figure 1. Posterior densities of selected coefficients

Sala-i-Martin (1997a,b). Two vertical dotted lines indicate a classical 90% confidence interval using the averaged variance of Sala-i-Martin (1997a,b), allowing for an informal comparison with his findings. The first two coefficients graphed in Figure 1 indicate the marginal effects on growth of two key variables: the initial level of GDP, and investment in equipment. The former was associated with conditional convergence of countries by Barro and Sala-i-Martin (1992), but this view was later challenged (Durlauf and Quah, 1999). In any case, it is widely accepted to be of theoretical and empirical importance and was one of the few regressors that Levine and Renelt (1992) found to be robust. The importance of Equipment investment was stressed in DeLong and Summers (1991). Whereas the inclusion of both variables receives strong support from BMA and Sala-i-Martin's classical analysis, the actual marginal inference on the regression coefficients is rather different, as can be judged from Figure 1. The next row of plots corresponds to two variables (Standard deviation of black market premium and Revolutions and coups) that are each included in Sala-i-Martin's analysis (with $CDF(0) \geq 0.993$), yet receive very low posterior probability of inclusion in our BMA analysis (3.1% and 1.7%, respectively). For the Standard deviation of the black market premium the confidence interval suggests much smaller spread than the (model averaged) posterior distribution. For Revolutions and coups we find that, in addition, the posterior mode is also quite far from the averaged classical point estimates. The same combination of differences in both spread and location is present for the next two coefficients: Absolute latitude and Spanish colony, which are both identified as important regressors in Sala-i-Martin, but get less than 3.5% posterior inclusion probability. The next two coefficients are illustrative of the opposite situation. The regressors Fraction Hindu and Black market premium are excluded in Sala-i-Martin (1997b), but BMA results indicate these are relatively important variables. Finally, Non-Equipment investment and Fraction Catholic are identified as important by Sala-i-Martin and also receive substantial posterior probability of inclusion in BMA. However, whereas the confidence interval accords reasonably well with the posterior results from BMA for Non-Equipment investment, it is quite different for Fraction Catholic: the second mode (with opposite sign!) is not at all picked up by the classical analysis.

Generally, averaged point estimates are often not too far from the posterior modes resulting from BMA, but most classical confidence intervals are far narrower than their posterior counterparts, thus severely underestimating uncertainty. This is not surprising in view of the fact that Sala-i-Martin's confidence intervals are based on the averaged variances (formula (5) in Sala-i-Martin, 1997b), which amounts to neglecting the (weighted) variance of point estimates across models. This also explains why $CDF(0)$ using approach 1 (as defined in footnote 1 and presented in Sala-i-Martin, 1997a) is virtually always larger than with approach 2.

4. PREDICTIVE RESULTS

An important quality of a model is that it can provide useful forecasts. In addition, such a predictive exercise immediately provides a benchmark for evaluating the model's adequacy. We consider predicting the observable y_f given the corresponding values of the regressors, grouped in a k -dimensional vector z_f (which has been transformed in the same way as Z —i.e. by subtracting the average of the original regressors over the n observations on which posterior inference is based—in order to assign the same interpretation to the regression coefficients in posterior and predictive analysis).

Prediction naturally fits in the Bayesian paradigm as all parameters can be integrated out, formally taking parameter uncertainty into account. If we also wish to deal with model uncertainty, BMA as in equation (5) provides us with the formal mechanism, and we can characterize the out-of-sample predictive distribution of y_f by

$$p(y_f|y) = \sum_{j=1}^{2^k} f_S \left(y_f \mid n-1, \bar{y} + \frac{1}{g+1} z'_{f,j} \beta_j^*, \right. \\ \left. \frac{n-1}{d_j^*} \times \left\{ 1 + \frac{1}{n} + \frac{1}{g+1} z'_{f,j} (Z'_j Z_j)^{-1} z_{f,j} \right\}^{-1} \right) P(M_j|y) \quad (9)$$

where $f_S(x|v, b, a)$ denotes the p.d.f. of a univariate Student- t distribution with v degrees of freedom, location b (the mean if $v > 1$) and precision a (with variance $v/\{a(v-2)\}$ provided $v > 2$) evaluated at x . In addition, $z_{f,j}$ groups the j elements of z_f corresponding to the regressors in M_j , $\beta_j^* = (Z'_j Z_j)^{-1} Z'_j y$ and

$$d_j^* = \frac{1}{g+1} y' M_{X_j} y + \frac{g}{g+1} (y - \bar{y} t_n)' (y - \bar{y} t_n) \quad (10)$$

We shall now split the sample into n observations on which we base our posterior inference and q observations which we retain in order to check the predictive accuracy of the model. As a formal criterion, we shall use the log predictive score (*LPS*), introduced by Good (1952). It is a strictly proper scoring rule, in the sense that it induces the forecaster to be honest in divulging his predictive distribution. For $f = n+1, \dots, n+q$ —i.e. for each country in the prediction sample—we base our measure on the predictive density evaluated in these retained observations y_{n+1}, \dots, y_{n+q} , namely:

$$LPS = -\frac{1}{q} \sum_{f=n+1}^{n+q} \ln p(y_f|y) \quad (11)$$

The smaller *LPS* is, the better the model does in forecasting the prediction sample. Interpreting values for *LPS* can perhaps be facilitated by considering that in the case of i.i.d. sampling, *LPS* approximates an integral that equals the sum of the Kullback–Leibler divergence between the actual sampling density and the predictive density in equation (9) and the entropy of the sampling distribution (Fernández *et al.*, 2001). So *LPS* captures uncertainty due to a lack of fit plus the inherent sampling uncertainty, and does not distinguish between these two. Here we are necessarily faced with a different z_f for every forecasted observation (corresponding to a specific country), so we are not strictly in the context of observations that are generated by the same distribution. Still, we think the above interpretation may shed some light on the calibration and comparison of *LPS* values. If, for the sake of argument, we assume that we fit the sampling distribution perfectly, then *LPS* approximates entropy alone. In the context of a Normal sampling model with fixed standard deviation σ_* , this latter entropy can then be translated into a value for σ_* , using the fact that entropy equals $\ln(\sigma_* \sqrt{2\pi e})$. Thus, a known Normal sampling distribution with fixed σ_* would induce the same inherent predictive uncertainty as measured by *LPS*, if we choose $\sigma_* = \exp(LPS)/\sqrt{2\pi e}$. Of course, as a direct consequence, a difference in *LPS* of, say, 0.1, corresponds to about a 10% difference in values for σ_* .

We shall use *LPS* to compare four different regression strategies: the BMA approach, leading to equation (9), the best model—i.e. the one with the highest posterior probability—in \mathcal{M} , the best

Table II. Predictive Performance

	Number of times			LPS		
	Best	Worst	Beaten by null	Min	Mean	Max
BMA	9	0	2	-3.470	-2.977	-2.408
Best model in \mathcal{M}	0	8	12	-3.370	-2.316	-1.268
Best model in \mathcal{M}_I	8	0	4	-3.460	-2.838	-1.341
Full model	3	8	12	-3.266	-2.261	-0.940
Null model	0	4	...	-2.850	-2.560	-1.853

model in \mathcal{M}_I , and the full model with all k regressors. As a benchmark for the importance of growth regression, we also include LPS for the ‘null model’, i.e. the model with only the intercept where no individual country characteristics are used. We would expect this model to reproduce the marginal growth distribution (without conditioning on regressors) pretty well. As the sample standard deviation of growth is 0.01813, we could then roughly expect LPS values for the null model around the entropy value corresponding to $\sigma_* = 0.01813$, which is -2.591 . For the regression models, we would hope that they predict better, since they use the information in the regressors, and this should ideally be reflected in a smaller LPS or a conditional σ_* value under 0.01813.

The partition of the sample into the inference and the prediction sample is done randomly, by assigning observations to the inference sample with probability 0.75. The results of twenty different partitions are summarized in Table II. Besides numerical summaries of the LPS values, we also indicate how often the model is beaten by the trivial null model, and how often the model performs best and worst. The following key characteristics emerge: the null model performs in a fairly conservative fashion and its mean LPS is very close to what we expected on the basis of the sample standard deviation. The full model and the best model in \mathcal{M} are beaten by the benchmark null model in over half the cases, and clearly perform worst of the regression models. The best model in \mathcal{M}_I does quite a bit better and generally improves a lot on the null model, but can sometimes lead to very bad predictions (the maximum value of LPS corresponds to $\sigma_* = 0.0633$, 3.5 times that of the sample). In contrast, the BMA approach never leads us far astray: it is only beaten by the null model twice and the largest value of LPS corresponds to $\sigma_* = 0.0218$ (and this occurs in a case where it actually outperforms all other models by a large margin). It performs best most frequently and the best prediction it produced corresponds to $\sigma_* = 0.0075$ which is only about 40% of the sample standard deviation. The mean LPS value for the BMA model corresponds to $\sigma_* = 0.0123$, i.e. a reduction of the sample standard deviation by about a third.

The fact that BMA does so much better than simply taking the best model is compelling evidence supporting the use of formal model averaging rather than the selection of any given model. Interestingly, the best model in \mathcal{M}_I does better than the best model in \mathcal{M} (which contains \mathcal{M}_I). This underlines that the highest posterior probability on the basis of the inference sample does not necessarily lead to the best predictions in the prediction sample. In addition, \mathcal{M}_I is restricted to those models that include the three regressors always retained by Sala-i-Martin on theory grounds. This extra information (although not always supported by the data) may help in predicting.¹¹

¹¹ Of course, this relative success of the best model in \mathcal{M}_I has no immediate bearing on the predictive performance of a classical analysis.

In summary, the use of regression models with BMA results in a considerable predictive improvement over the null model, thus clearly suggesting that growth regression is not a futile exercise, although care should be taken in the methodology adopted.

5. DISCUSSION

The value of growth regression in cross-country analysis has been illustrated in the predictive exercise in the previous section. We agree with Sala-i-Martin (1997b) that some regressors can be identified as useful explanatory variables for growth in a linear regression model, but we advocate a formal treatment of model (and parameter) uncertainty. In our methodology the marginal importance of an explanatory variable does not necessarily imply anything about the size or sign of the regression coefficient in a set of models, but is based entirely on the posterior probabilities of models containing that regressor. In addition, we go one step further and provide a practical and theoretically sound method for inference, both posterior and predictive, using Bayesian Model Averaging (BMA). From the huge spread of the posterior mass in model space and the predictive advantage of BMA, it is clear that model averaging is recommended when dealing with growth regression.

Our Bayesian paradigm provides us with a formal framework to implement this model averaging, and recent Markov chain Monte Carlo (MCMC) methods are shown to be very powerful indeed. Despite the huge model space (with 2.2 trillion models), we obtain very reliable results without an inordinate amount of computational effort.¹²

The analysis in Sala-i-Martin (1997b) is not Bayesian and thus no formal model averaging can occur, even though he considers weighing with the integrated likelihood.¹³ In addition, the latter analysis evaluates all models and is thus necessarily restricted to a rather small set of models, \mathcal{M}_S , which seems not to receive that much empirical support from the data. Even though we find a roughly similar set of variables that can be classified as 'important' for growth regressions, a crucial additional advantage is that our results are immediately interpretable in terms of model probabilities and all inference can easily be conducted in a purely formal fashion by BMA. It is not clear to us what to make of the recommendations in Sala-i-Martin (1997b): should the applied researcher use all the regressors identified as important or mix over the corresponding models in \mathcal{M}_S ? However, the latter would have to be without proper theoretical foundation or guidance if a classical statistical framework is adopted.

In our view, the treatment of a very large model set, such as \mathcal{M} , in a theoretically sound and empirically practical fashion requires BMA and MCMC methods. In addition, this methodology provides a clear and precise interpretation of the results, and immediately leads to posterior and predictive inference.

ACKNOWLEDGEMENTS

We greatly benefited from useful comments from co-editor Steven Durlauf, Ed Leamer, Jon Temple, an anonymous referee, and participants at seminars at the Fundación de Estudios de Economía Aplicada (Madrid), the Inter-American Development Bank, the International Monetary

¹² The reported chain took about $2\frac{2}{3}$ hours of CPU time on a fast Sun Ultra-2 with two 296MHz CPUs, 512Mb of RAM and 3.0Gb of swap space running under Solaris 2.6. Our programs are coded in Fortran 77 and are available at this journal's website.

¹³ Sala-i-Martin (1997a,b) do not specify what this integrated likelihood is; as there is no prior to integrate with, this may refer to the maximized likelihood, which is proportional to $(y'M_{X_j}y)^{-n/2}$ for M_j .

Fund, Heriot-Watt University, the University of Kent at Canterbury, the University of Southampton, and the 1999 European Economic Association Meetings at Santiago de Compostela. Part of this research was conducted while Carmen Fernández was at the Department of Mathematics, University of Bristol, and Mark Steel at the Department of Economics, University of Edinburgh.

REFERENCES

- Barro RJ. 1991. Economic growth in a cross section of countries. *Quarterly Journal of Economics* **106**: 407–444.
- Barro RJ, Sala-i-Martin XX. 1992. Convergence. *Journal of Political Economy* **100**: 223–251.
- Brock WA, Durlauf SN. 2000. Growth economics and reality. Working Paper 8041, National Bureau of Economic Research, Cambridge, Massachusetts.
- Chib S, Greenberg E. 1996. Markov chain Monte Carlo simulation methods in econometrics. *Econometric Theory* **12**: 409–431.
- Chipman H. 1996. Bayesian variable selection with related predictors. *Canadian Journal of Statistics* **24**: 27–36.
- Clyde M, Desimone H, Parmigiani G. 1996. Prediction via orthogonalized model mixing. *Journal of the American Statistical Association* **91**: 1197–1208.
- DeLong JB, Summers L. 1991. Equipment investment and economic growth. *Quarterly Journal of Economics* **106**: 445–502.
- Durlauf SN, Quah DT. 1999. The new empirics of economic growth. In *Handbook of Macroeconomics* (Vol. 1A), Taylor JB, Woodford M (eds). North-Holland: Amsterdam; 231–304.
- Fernández C, Ley E, Steel MFJ. 2001. Benchmark priors for Bayesian model averaging. *Journal of Econometrics* **100**: 381–427.
- George EI. 1999. Bayesian model selection. In *Encyclopedia of Statistical Sciences Update* (Vol. 3), Kotz S, Read C, Banks DL (eds). Wiley: New York; 39–46.
- George EI, McCulloch RE. 1997. Approaches for Bayesian variable selection. *Statistica Sinica* **7**: 339–373.
- Good IJ. 1952. Rational decisions. *Journal of the Royal Statistical Society, B* **14**: 107–114.
- Kass RE, Raftery AE. 1995. Bayes factors. *Journal of the American Statistical Association* **90**: 773–795.
- Kormendi R, Meguire P. 1985. Macroeconomic determinants of growth, cross-country evidence. *Journal of Monetary Economics* **16**: 141–163.
- Leamer EE. 1978. *Specification Searches: Ad Hoc Inference with Nonexperimental Data*. Wiley: New York.
- Leamer EE. 1983. Let's take the con out of econometrics. *American Economic Review* **73**: 31–43.
- Leamer EE. 1985. Sensitivity analyses would help. *American Economic Review* **75**: 308–313.
- Lee H. 1996. Model selection for consumer loan application data. Dept of Statistics Working Paper No. 650, Carnegie-Mellon University.
- Levine R, Renelt D. 1992. A sensitivity analysis of cross-country growth regressions. *American Economic Review* **82**: 942–963.
- Madigan D, York J. 1995. Bayesian graphical models for discrete data. *International Statistical Review* **63**: 215–232.
- Mitchell TJ, Beauchamp JJ. 1988. Bayesian variable selection in linear regression. *Journal of the American Statistical Association* **83**: 1023–1036 (with discussion).
- Raftery AE, Madigan D, Hoeting JA. 1997. Bayesian model averaging for linear regression models. *Journal of the American Statistical Association* **92**: 179–191.
- Sala-i-Martin XX. 1997a. I just ran four million regressions. Mimeo, Columbia University.
- Sala-i-Martin XX. 1997b. I just ran two million regressions. *American Economic Review* **87**: 178–183.
- Smith M, Kohn R. 1996. Nonparametric regression using Bayesian variable selection. *Journal of Econometrics* **75**: 317–343.
- Temple J. 1999. The new growth evidence. *Journal of Economic Literature* **37**: 112–156.
- Temple J. 2000. Growth regressions and what the textbooks don't tell you. *Bulletin of Economic Research* **52**: 181–205.
- Tierney L. 1994. Markov chains for exploring posterior distributions. *Annals of Statistics* **22**: 1701–1762 (with discussion).